

In [1]:

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import collections
import seaborn as sns
%matplotlib inline
```

In [2]:

```
ratings_data = pd.read_csv('ratings.csv')
movies_df = pd.read_csv("movies.csv")
tags = pd.read_csv('tags.csv')
scores = pd.read_csv('genome-scores.csv')
links = pd.read_csv('links.csv')
genome_tags = pd.read_csv('genome-tags.csv')
```

In [3]:

```
movies_df.head()
```

Out[3]:

| | movieId | title | genres |
|---|---------|------------------------------------|---|
| 0 | 1 | Toy Story (1995) | Adventure Animation Children Comedy Fantasy |
| 1 | 2 | Jumanji (1995) | Adventure Children Fantasy |
| 2 | 3 | Grumpier Old Men (1995) | Comedy Romance |
| 3 | 4 | Waiting to Exhale (1995) | Comedy Drama Romance |
| 4 | 5 | Father of the Bride Part II (1995) | Comedy |

In [40]:

```
movies_df.shape
```

Out[40]:

```
(27278, 3)
```

In [42]:

```
movies_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27278 entries, 0 to 27277
Data columns (total 3 columns):
movieId    27278 non-null int64
title      27278 non-null object
genres     27278 non-null object
dtypes: int64(1), object(2)
memory usage: 639.5+ KB
```

In [44]:

```
movies_df.isnull().sum()
```

Out[44]:

```
movieId    0
title      0
genres     0
dtype: int64
```

In [41]:

```
movies_df.describe()
```

Out[41]:

| | movieId |
|-------|---------------|
| count | 27278.000000 |
| mean | 59855.480570 |
| std | 44429.314697 |
| min | 1.000000 |
| 25% | 6931.250000 |
| 50% | 68068.000000 |
| 75% | 100293.250000 |
| max | 131262.000000 |

In [4]:

```
movie_genres = []

for genre in movies_df['genres']:
    for movie in genre.split('|'):
        movie_genres.append(movie)
```

In [5]:


```
genre_counts = pd.Series(movie_genres).value_counts()[:18]
```

In [6]:

```
top_genres = pd.Series(movie_genres).value_counts()[:10]

sns.barplot(y=top_genres.index, x=top_genres.values, palette='magma').set_title(
    'Top-10 Movie Genres', fontsize=14, weight='bold')

plt.show()
```



Drama is the most commonly occurring genre with almost half the movies identifying itself as a drama film. Comedy comes in at a distant second with 30% of the movies having adequate doses of humor. Other major genres represented in the top 10 are Thriller, Romance, Action, Crime, Horror, Documentry, Adventure, SciFi.

We will consider only those themes that appear in the top 10 most popular genres.

In [7]:

```
ratings_data.head()
```

Out[7]:

Out[7]:

| | userId | movieId | rating | timestamp |
|---|--------|---------|--------|------------|
| 0 | 1 | 2 | 3.5 | 1112486027 |
| 1 | 1 | 29 | 3.5 | 1112484676 |
| 2 | 1 | 32 | 3.5 | 1112484819 |
| 3 | 1 | 47 | 3.5 | 1112484727 |
| 4 | 1 | 50 | 3.5 | 1112484580 |

In [8]:

```
ratings_data['timestamp'] = pd.to_datetime(ratings_data['timestamp'], unit='s')
ratings_data.head()
```

Out[8]:

| | userId | movieId | rating | timestamp |
|---|--------|---------|--------|---------------------|
| 0 | 1 | 2 | 3.5 | 2005-04-02 23:53:47 |
| 1 | 1 | 29 | 3.5 | 2005-04-02 23:31:16 |
| 2 | 1 | 32 | 3.5 | 2005-04-02 23:33:39 |
| 3 | 1 | 47 | 3.5 | 2005-04-02 23:32:07 |
| 4 | 1 | 50 | 3.5 | 2005-04-02 23:29:40 |

In [9]:

```
ratings_data.shape
```

Out[9]:

(20000263, 4)

In [10]:

```
ratings_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000263 entries, 0 to 20000262
Data columns (total 4 columns):
userId      int64
movieId     int64
rating      float64
timestamp   datetime64[ns]
dtypes: datetime64[ns](1), float64(1), int64(2)
memory usage: 610.4 MB
```

In [11]:

```
ratings_data.describe()
```

Out[11]:

| | userId | movieId | rating |
|-------|--------------|--------------|--------------|
| count | 2.000026e+07 | 2.000026e+07 | 2.000026e+07 |
| mean | 6.904587e+04 | 9.041567e+03 | 3.525529e+00 |
| std | 4.003863e+04 | 1.978948e+04 | 1.051989e+00 |
| min | 1.000000e+00 | 1.000000e+00 | 5.000000e-01 |
| 25% | 3.439500e+04 | 9.020000e+02 | 3.000000e+00 |
| 50% | 6.914100e+04 | 2.167000e+03 | 3.500000e+00 |
| 75% | 1.036370e+05 | 4.770000e+03 | 4.000000e+00 |
| max | 1.381930e+05 | 1.312620e+05 | 5.000000e+00 |

```
max 1.000000e+00 1.012020e+00 0.000000e+00
      userid      movied      rating
```

In [12]:

```
ratings_data.corr()
```

Out[12]:

| | userid | movied | rating |
|--------|-----------|-----------|----------|
| userid | 1.000000 | -0.000638 | 0.001175 |
| movied | -0.000638 | 1.000000 | 0.001191 |
| rating | 0.001175 | 0.001191 | 1.000000 |

In [13]:

```
ratings_data['gave_rating_year'] = ratings_data['timestamp'].dt.year
ratings_data['gave_rating_month'] = ratings_data['timestamp'].dt.month_name().str[:3]

ratings_data.drop('timestamp', axis=1, inplace=True)
```

In [14]:


```
ratings_data.head()
```

Out[14]:

| | userid | movied | rating | gave_rating_year | gave_rating_month |
|---|--------|--------|--------|------------------|-------------------|
| 0 | 1 | 2 | 3.5 | 2005 | Apr |
| 1 | 1 | 29 | 3.5 | 2005 | Apr |
| 2 | 1 | 32 | 3.5 | 2005 | Apr |
| 3 | 1 | 47 | 3.5 | 2005 | Apr |
| 4 | 1 | 50 | 3.5 | 2005 | Apr |

In [15]:

```
ratings_data['rating'].value_counts().plot(kind='bar');
```



In [16]:

```
ratings_data.rating.value_counts(normalize=True)
```

Out[16]:

```
4.0    0.278093
3.0    0.214557
5.0    0.144931
3.5    0.110006
4.5    0.076740
2.0    0.072510
2.5    0.047510
1.0    0.035010
1.5    0.015010
0.5    0.010010
```

```
2.0    0.071549
2.5    0.044169
1.0    0.034036
1.5    0.013962
0.5    0.011956
Name: rating, dtype: float64
```

In [45]:

```
ratings_data.isnull().sum()
```

Out[45]:

```
userId      0
movieId     0
rating      0
gave_rating_year  0
gave_rating_month 0
dtype: int64
```

In [17]:

```
movies_with_ratings = ratings_data.merge(movies_df, on='movieId', how='inner')
movies_with_ratings.head()
```

Out[17]:

| | userId | movieId | rating | gave_rating_year | gave_rating_month | title | genres |
|---|--------|---------|--------|------------------|-------------------|----------------|----------------------------|
| 0 | 1 | 2 | 3.5 | 2005 | Apr | Jumanji (1995) | Adventure Children Fantasy |
| 1 | 5 | 2 | 3.0 | 1996 | Dec | Jumanji (1995) | Adventure Children Fantasy |
| 2 | 13 | 2 | 3.0 | 1996 | Nov | Jumanji (1995) | Adventure Children Fantasy |
| 3 | 29 | 2 | 3.0 | 1996 | Jun | Jumanji (1995) | Adventure Children Fantasy |
| 4 | 34 | 2 | 3.0 | 1996 | Oct | Jumanji (1995) | Adventure Children Fantasy |

In [18]:

```
counts = movies_with_ratings['title'].value_counts()[:10]
counts
```

Out[18]:


```
Pulp Fiction (1994)          67310
Forrest Gump (1994)          66172
Shawshank Redemption, The (1994) 63366
Silence of the Lambs, The (1991) 63299
Jurassic Park (1993)         59715
Star Wars: Episode IV - A New Hope (1977) 54502
Braveheart (1995)            53769
Terminator 2: Judgment Day (1991) 52244
Matrix, The (1999)           51334
Schindler's List (1993)      50054
Name: title, dtype: int64
```


In [19]:

```
top_genres = pd.Series(counts)

sns.barplot(y=top_genres.index, x=top_genres.values, palette='magma').set_title(
    'Top-10 popular Movies', fontsize=14, weight='bold', color = 'skyblue')

plt.show()
```






In [20]:

```
ratings_data.boxplot(column='rating', figsize=(10,5))
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x1eeb92c1630>



Basic Statistics (#Ratings, #Users, and #Movies)

In [21]:

```
print("Total data ")
print("-"*50)
print("\nTotal no of ratings :", ratings_data.shape[0])
print("Total No of Users   :", len(np.unique(ratings_data.userId)))
print("Total No of movies  :", len(np.unique(ratings_data.movieId)))
```


Total data


```
-----
Total no of ratings : 20000263
Total No of Users   : 138493
Total No of movies  : 26744
```

In [22]:

```
fig, ax = plt.subplots()
plt.title('Distribution of ratings over Training dataset', fontsize=15)
sns.countplot(ratings_data.rating)
#ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
ax.set_ylabel('No. of Ratings(Millions)')

plt.show()
```





In [23]:

```
no_of Rated movies per user = ratings_data.groupby (by='userId')
['rating'].count().sort_values(ascending=False)

no_of Rated movies per user.head()
```

Out[23]:

```
userId
118205    9254
8405      7515
82418     5646
121535    5520
125794     5491
Name: rating, dtype: int64
```

In [24]:

```
no_of Rated movies per user.describe()
```

Out[24]:

```
count    138493.000000
mean      144.413530
std       230.267257
min        20.000000
25%       35.000000
50%       68.000000
75%      155.000000
max      9254.000000
Name: rating, dtype: float64
```

In [25]:

```
no_of Rated movies per user.min()
```

Out[25]:

```
20
```

In [26]:

```
no_of Rated movies per user.max()
```

Out[26]:

```
9254
```

In [27]:


```
fig = plt.figure(figsize=plt.figaspect(.5))

ax1 = plt.subplot(121)
sns.kdeplot(no_of Rated movies per user, shade=True, ax=ax1)
plt.xlabel('No of ratings by user')
plt.title("PDF")

ax2 = plt.subplot(122)
sns.kdeplot(no_of Rated movies per user, shade=True, cumulative=True, ax=ax2)
plt.xlabel('No of ratings by user')
```

```
plt.xlabel('No of ratings by user')
plt.title('CDF')

plt.show()
```

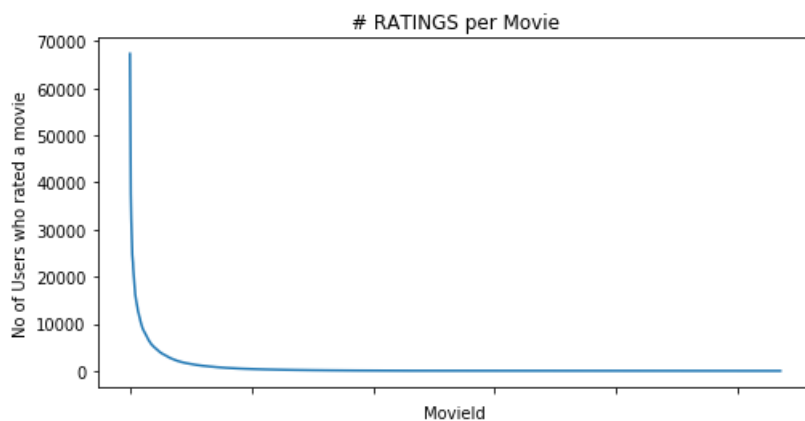


In [28]:

```
no_of_ratings_per_movie = ratings_data.groupby(by='movieId')
['rating'].count().sort_values(ascending=False)

fig = plt.figure(figsize=plt.figaspect(.5))
ax = plt.gca()
plt.plot(no_of_ratings_per_movie.values)
plt.title('# RATINGS per Movie')
plt.xlabel('MovieId')
plt.ylabel('No of Users who rated a movie')
ax.set_xticklabels([])

plt.show()
```



It is very skewed.. just like number of ratings given per user.

- There are some movies (which are very popular) which are rated by huge number of users.
- But most of the movies(like 90%) got some hundreds of ratings.

In [29]:

```
no_of Rated movies per user.describe()
```

Out[29]:

```
count    138493.000000
mean       144.413530
std        230.267257
min         20.000000
25%         35.000000
50%         68.000000
75%        155.000000
max        9254.000000
```



```
max      9254.000000
Name: rating, dtype: float64
```

In [30]:


```
quantiles = no_of_rated_movies_per_user.quantile(np.arange(0,1.01,0.01), interpolation='higher')
```

In [31]:

```
plt.title("Quantiles and their Values")
quantiles.plot()
# quantiles with 0.05 difference
plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label="quantiles with 0.05 intervals")
# quantiles with 0.25 difference
plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label = "quantiles with 0.25 intervals")
plt.ylabel('No of ratings by user')
plt.xlabel('Value at the quantile')
plt.legend(loc='best')

# annotate the 25th, 50th, 75th and 100th percentile values....
for x,y in zip(quantiles.index[::25], quantiles[::25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
                , fontweight='bold')

plt.show()
```



In [32]:

```
quantiles[::5]
```

Out [32]:

```
0.00    20
0.05    21
0.10    24
0.15    27
0.20    30
0.25    35
0.30    39
0.35    45
0.40    51
0.45    59
0.50    68
0.55    79
0.60    93
0.65   108
0.70   127
0.75   155
0.80   193
0.85   246
0.90   334
0.95   520
1.00  9254
```

Name: rating, dtype: int64

how many ratings at the last 5% of all ratings??

In [33]:

```
print('\n No of ratings at last 5 percentile : {} \n'.format(sum(no_of Rated_movies_per_user >= 520)
) )
```

No of ratings at last 5 percentile : 6940

In [34]:


```
ratings_data.head(3)
```

Out[34]:

| | userid | movieId | rating | gave_rating_year | gave_rating_month |
|---|--------|---------|--------|------------------|-------------------|
| 0 | 1 | 2 | 3.5 | 2005 | Apr |
| 1 | 1 | 29 | 3.5 | 2005 | Apr |
| 2 | 1 | 32 | 3.5 | 2005 | Apr |


In [35]:


```
fig, ax = plt.subplots()
sns.countplot(x='gave_rating_month', data=ratings_data, ax=ax)
plt.title('No of ratings on each month...')
plt.ylabel('Total no of ratings')
plt.xlabel('')
#ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```



In [36]:

```
from datetime import datetime
start = datetime.now()
fig = plt.figure(figsize=plt.figaspect(.45))
sns.boxplot(y='rating', x='gave_rating_month', data=ratings_data)
plt.show()
print(datetime.now() - start)
```





0:00:08.468619

In [37]:

```
avg_month_df = ratings_data.groupby(by=['gave_rating_month'])['rating'].mean()
print(" AVerage ratings")
print("-"*30)
print(avg_month_df)
print("\n")
```

```
 AVerage ratings
-----
gave_rating_month
Apr      3.528203
Aug      3.505279
Dec      3.539704
Feb      3.524092
Jan      3.520437
Jul      3.517270
Jun      3.513506
Mar      3.496393
May      3.514545
Nov      3.546288
Oct      3.565256
Sep      3.528368
Name: rating, dtype: float64
```

In []: